

XI'AN JIAOTONG-LIVERPOOL UNIVERSITY

西 交 利 物 浦 大 学

COURSEWORK SUBMISSION COVER SHEET

Name	Changqing Lin	(Other Names)
Student Number	2039153	
Programme	BSc Information Management and Information Systems	
Module Title	Final Year Project	
Module Code	BUS303	
Assignment Title	FYP Dissertation	
Submission Deadline	9:00 on 13 th May 2024	
Module Leader	Dr. Yidong Tao	

By uploading or submitting this coursework submission cover sheet, I certify the following:

- I have read and understood the definitions of collusion, copying, plagiarism, and dishonest use of data as outlined in the Academic Integrity Policy of Xi'an Jiaotong-Liverpool University.
- This work is my own, original work produced specifically for this assignment. It does not misrepresent the work of another person or institution as my own. Additionally, it is a submission that has not been previously published, or submitted to another module
- This work is not the product of unauthorized collaboration between myself and others.
- This work is free of embellished or fabricated data.

I understand collusion, plagiarism, dishonest use of data, and submission of procured work are serious academic misconducts. By uploading or submitting this cover sheet, I acknowledge that I am subject to disciplinary action if I am found to have committed such acts.

SignatureChangqing Lin.....

Date2024/5/12.....

For Academic Office use:	Date Received	Working Days Late	Penalty ¹

¹ Please refer to clause 43 (i), (ii) and (iii) of the Code of Practice on Assessment for the standard system of penalties for the late submission of assessment work based on **working days**.

**Predicting Future Performance of Financial Assets Using Market
Indicators: A Machine Learning and Neural Network Approach**

By

Changqing Lin

2039153

Supervisor: Dr. Jiatao Liu

**A Dissertation Submitted in Partial Fulfillment of the Requirements for the
BSc Information Management and Information Systems**

to

International Business School Suzhou

Xi'an Jiaotong-Liverpool University

May 13, 2024

Abstract

This research aims to enhance the accuracy of stock price predictions by leveraging advanced neural network architectures, specifically CNN-MLP, AlexNet-MLP, and ResNet-MLP. The study explores the effectiveness of these hybrid models in capturing complex financial data patterns and compares their performance using different optimizers. The research employs a quantitative methodology, utilizing historical stock price data from the S&P 500 index. The models are trained and evaluated using metrics such as RMSE, MAE, R-squared, Total Return, and Annualized Return. Various learning rates and optimizers (Adam and SGD) are tested to determine optimal configurations. The findings indicate that models optimized with the Adam optimizer generally outperform those using SGD, due to Adam's adaptive learning rate. The CNN-MLP model with Adam demonstrates the best overall performance in terms of accuracy and profitability. The AlexNet-MLP and ResNet-MLP models also show competitive performance, particularly in handling complex data patterns.

Keywords: Price Prediction, Neural Networks, CNN, AlexNet, ResNet, Adam Optimizer, Experiment

摘要

本研究旨在通过利用先进的神经网络架构，特别是CNN-MLP、AlexNet-MLP和ResNet-MLP，提高股票价格预测的准确性。研究探索了这些混合模型在捕捉复杂金融数据模式中的有效性，并比较了它们在不同优化器下的表现。在设计上本研究采用定量方法，利用标普500指数的历史股票价格数据。模型通过RMSE、MAE、R平方、总回报率和年化回报率等指标进行训练和评估。测试了各种学习率和优化器（Adam和SGD）以确定最佳配置。研究结果表明，由于Adam的自适应学习率，使用Adam优化器的模型通常优于使用SGD的模型。CNN-MLP模型在准确性和盈利能力方面表现最佳。AlexNet-MLP和ResNet-MLP模型在处理复杂数据模式方面也表现出色。

关键词：价格预测, 神经网络, 卷积神经网络, AlexNet, 残差神经网络, Adam优化器, 实验

Acknowledgement

I would like to extend my deepest gratitude to my supervisor, Dr. Jiatao Liu, for his unwavering support, invaluable guidance, and continuous encouragement throughout the course of this project. His insightful feedback and expert advice were instrumental in shaping this research, and his constant encouragement inspired me to strive for excellence. I am profoundly grateful for his mentorship, which has been a cornerstone in the successful completion of this dissertation. Thank you, Dr. Liu, for your dedication and support.

Table of Contents

COURSEWORK COVERSHEET	I
TITLE PAGE.....	II
ABSTRACT	III
ACKNOWLEDGEMENT	IV
TABLE OF CONTENTS	V
LIST OF TABLES.....	VII
LIST OF FIGURES.....	VII
LIST OF ACRONYMS.....	VII
1. INTRODUCTION.....	1
1.1 BACKGROUND.....	1
1.2 MOTIVATION.....	1
1.3 DIFFERENTIATION	2
1.4 RESEARCH METHOD.....	2
1.5 PREVIEW OF RESEARCH RESULTS	2
1.6 CONTRIBUTION	3
2. LITERATURE REVIEW.....	3
3. METHODOLOGY.....	7
3.1 MODEL ARCHITECTURES	7
3.1.1 CNN.....	7
3.1.2 AlexNet	9
3.1.3 ResNet.....	10
3.1.4 MLP.....	10
3.2 OPTIMIZATION METHODS.....	12
3.2.1 SGD.....	12
3.2.2 Adam.....	13
3.3 EXPERIMENTAL DESIGN	14
3.3.1 <i>Experimental Setting, Benefits, and Influence</i>	14
3.3.2 <i>Stock Market Data Collection</i>	15
3.3.3 <i>Data Preprocessing</i>	15
3.3.4 <i>Model Code Design</i>	16
3.3.5 <i>CNN-MLP Model</i>	17
3.3.6 <i>CNN&AlexNet&ResNet-MLP Training and Prediction Process</i>	18
3.3.7 <i>Training Procedure</i>	19
3.3.8 <i>Trading Strategy Simulation</i>	20
4. EMPIRICAL RESULT.....	21
4.1 PERFORMANCE COMPARISON BASED ON MODEL ARCHITECTURES AND OPTIMIZERS	21
4.1.1 <i>CNN-MLP Performance</i>	22
4.1.2 <i>CNN(AlexNet)-MLP Performance</i>	22
4.1.3 <i>CNN(ResNet)-MLP Performance</i>	23
4.2 PERFORMANCE OF KEY METRICS	24
4.2.1 <i>RMSE and MAE Analysis</i>	24
4.2.2 <i>Total Return and Annualized Return Analysis</i>	24
4.3 SUMMARY OF FINDINGS	25
5. DISCUSSION	25
5.1 INTERPRETATION OF RESULTS.....	25

5.2 COMPARISON WITH LITERATURE.....	26
5.3 INSIGHTS FOR PRACTICE.....	28
6. CONCLUSION.....	28
REFERENCES	30
APPENDIX: ALL USED DATA AND CODE.....	35

List of Tables

TABLE 1: CNN-MLP PERFORMANCE	22
TABLE 2: CNN(ALEXNET)-MLP PERFORMANCE	23
TABLE 3: CNN(RESNET)-MLP PERFORMANCE.....	23
TABLE 4: RMSE & MAE PERFORMANCE	24
TABLE 5: RETURN METRICS COMPARISON	24

List of Figures

FIGURE 1: THE PROCESS OF CNN-MLP MODEL	19
FIGURE 2. CNN(ADAM)-MLP PREDICTION RESULTS	27
FIGURE 3. CNN(ALEXNET)-MLP PREDICTION RESULTS	27
FIGURE 4. CNN(RESNET)-MLP PREDICTION RESULTS	28
FIGURE 5: EXPERIMENTAL RESULTS COLLECTION	37

List of Acronyms

Term	Initial components of the term
CNN	Convolutional Neural Network
SGD	Stochastic Gradient Descent
Adam	Adaptive Moment Estimation
MLP	Multilayer Perceptron
ResNet	Residual Neural Network

1. Introduction

1.1 Background

The intricate fluctuations of stock prices represent a complex puzzle in the economic domain, drawing significant interest from researchers and practitioners alike (Vanaga & Sloka, 2020). Influenced by a myriad of factors, both internal and external, stock prices are a reflection of various elements including the domestic and international economic environment, political scenarios, industry prospects, and the operational aspects of the stock market (Zhang & Kim, 2020; Badea, Ionescu & Guzun, 2019).

Historically, the primary methodology for stock price prediction has been rooted in economics and finance, utilizing two main approaches: fundamental analysis and technical analysis. Fundamental analysis delves into the intrinsic value of stocks, examining external influences like interest rates, exchange rates, inflation, industrial policies, financial health of listed companies, and geopolitical factors (Sousa, Montevechi & Miranda, 2019). On the other hand, technical analysis concentrates on the directional movement of stock prices, trading volumes, and investor psychology, employing tools like K-line charts to analyze the trajectory of stock indices (Coser, Maer-Matei & Albu, 2019).

However, the reliability of traditional fundamental analysis methods in predicting stock prices has been a subject of debate. The limitations stem not only from the long-term cycles of influencing factors but also from the dependency on the analysts' expertise. Given the random walk nature of financial time series, some scholars have turned to statistical and probabilistic models, such as the vector autoregression (VAR) (Jung & Boyd, 1996), Bayesian vector autoregression (BVAR) (Bleesser & Liicoff, 2005), autoregressive integrated moving average (ARIMA) (Adebiyi, Adewumi & Ayo, 2014), and generalized autoregressive conditional heteroskedasticity (GARCH) (Zhang, Cheng & Wang, 2005), to forecast short-term stock prices. Yet, the accuracy of these models is often questioned due to the inherent uncertainty and noise in financial data, alongside the dynamic nature of the relationships between variables (Yang & Wang, 2019).

1.2 Motivation

The main motivation for this research is the growing recognition of the limitations of traditional stock price prediction methods and the potential of advanced neural network architectures to provide more accurate and robust predictions. The significant volatility and complexity of financial markets necessitate sophisticated analytical tools that can capture intricate patterns and dependencies within the data. This research aims to explore and enhance the predictive capabilities of hybrid neural network models, particularly CNN-MLP, AlexNet-MLP, and ResNet-MLP, to offer more reliable and actionable insights for investors and financial analysts.

1.3 Differentiation

While existing literature has explored various machine learning and statistical models for stock price prediction, there is a notable gap in the application and comparative analysis of advanced neural network architectures like AlexNet and ResNet in this context. This research distinguishes itself by adapting these architectures, traditionally used in image recognition, for financial time series prediction. This novel approach seeks to evaluate their effectiveness relative to the more conventional CNN-MLP model, thereby potentially shedding new light on the advantages of deeper, more complex neural networks in financial forecasting.

1.4 Research Method

The research employs a quantitative methodology, implementing CNN-MLP, AlexNet-MLP, and ResNet-MLP models to predict stock prices. The models are trained and tested on historical stock price data, with performance assessed using metrics such as Root Mean Squared Error (RMSE), Mean Absolute Error (MAE), R-squared (R^2), Total Return, and Annualized Return. The study systematically varies the learning rates, optimizer functions, and proportions of training data to identify the optimal configurations for each model.

1.5 Preview of Research Results

Preliminary findings indicate that models optimized with the Adam optimizer generally outperform those using SGD, thanks to Adam's adaptive learning rate, which effectively handles the sparse gradients often present in financial data. Among the models, ResNet-MLP demonstrates superior performance in capturing complex patterns and dependencies, likely due to its residual learning framework, which mitigates the vanishing gradient problem and allows for deeper network training.

1.6 Contribution

This research contributes to the field of financial analytics by providing empirical evidence on the efficacy of advanced neural network architectures for stock price prediction. The findings suggest that integrating more complex models like AlexNet and ResNet can enhance predictive accuracy and financial returns, offering valuable insights for both algorithm developers and financial practitioners. The study also underscores the importance of selecting appropriate optimizers and tuning model parameters to maximize predictive performance and profitability.

2. Literature Review

A large amount of research has been conducted in the field of stock price forecasting, but there is not yet a unified methodology that can accurately predict the future behavior of the market. Traditional forecasting methods such as fundamental and technical analysis provide some level of understanding of market behaviors, but these methods usually fail to capture non-linear patterns in market data (Smith & Jones, 2015; Taylor, 2018). For the analysis of financial time series, classical statistical models such as VAR and GARCH have been widely used for short-term stock price forecasting (Brown & Miao, 2019; Liu et al., 2020).

In recent years, with the rapid development of machine learning and deep learning techniques, researchers have begun to explore the application of these methods in stock price prediction. In particular, the combination of CNNs and MLPs has shown significant advantages in capturing spatial and temporal features of time series data (Johnson & Zhang, 2017; Kumar & Patel, 2019). These studies have shown that hybrid models can handle large amounts of historical data more efficiently, thereby improving forecasting accuracy.

Convolutional neural networks (CNNs) have revolutionized the field of pattern recognition since their inception by LeCun et al. (1998). Their architecture, designed to process grid-like data structures, excels in feature extraction through a series of convolutional layers that capture hierarchical patterns within the data (LeCun et al., 1998). This capability has been leveraged in various domains, including financial time series analysis, where capturing temporal dependencies and intricate patterns is crucial for accurate forecasting (Kim & Kim, 2019). CNNs' ability to reduce the

dimensionality of the data while preserving its essential features makes them particularly effective in handling the high-dimensional nature of financial datasets.

Multilayer perceptrons (MLPs), characterized by their fully connected layers, have been pivotal in the evolution of artificial neural networks. MLPs are adept at modeling complex, non-linear relationships between input features and target variables, making them particularly effective in handling multi-dimensional data (Glorot & Bengio, 2010). Each neuron in an MLP employs a non-linear activation function on the weighted sum of its inputs, facilitating the network's ability to learn and represent intricate patterns in the data (Nair & Hinton, 2010). This non-linearity is essential for capturing the complex dynamics of financial time series, where relationships between variables are often non-linear and multifaceted. The flexibility of MLPs enables their integration with other neural network architectures, enhancing their capability to learn from diverse data types and improving the model's overall predictive performance (Goodfellow, Bengio & Courville, 2016).

AlexNet, introduced by Krizhevsky, Sutskever, and Hinton (2012), marked a significant advancement in deep learning. Its deep architecture, combined with innovative use of ReLU activation functions and dropout for regularization, has set new benchmarks in image classification tasks. AlexNet's robust feature extraction capabilities through deep convolutional and pooling layers make it highly suitable for handling high-dimensional financial data, allowing for enhanced pattern recognition and trend analysis in stock price movements (Krizhevsky et al., 2012). The architecture's success in large-scale image classification tasks suggests its potential effectiveness in financial forecasting, where capturing detailed temporal and spatial patterns is critical.

ResNet, developed by He et al. (2016), addressed the vanishing gradient problem prevalent in very deep networks through the introduction of residual learning. By employing identity mappings and shortcut connections, ResNet preserves information flow across layers, enhancing feature extraction capabilities and enabling the training of much deeper networks. This architecture is particularly advantageous for financial forecasting, where capturing long-term dependencies and complex temporal patterns is essential (He et al., 2016). The residual learning approach of ResNet ensures that essential features are retained across layers, making it effective in scenarios where detailed pattern recognition is necessary.

The application of CNNs in financial time series forecasting has been explored in various studies. For instance, Heaton, Polson, and Witte (2016) demonstrated the effectiveness of CNNs in capturing temporal dependencies and patterns in financial data, leading to improved predictive performance. They highlighted CNNs' ability to extract meaningful features from raw time series data, facilitating more accurate predictions. Similarly, Qin, Yu, and Zhao (2018) applied a deep learning approach to forecast stock prices in the Chinese market, highlighting the advantages of CNNs in financial contexts. Their study demonstrated that CNNs could effectively capture and utilize the temporal and spatial dependencies present in financial data.

MLPs have also been extensively researched in financial forecasting. Kamalov (2020) utilized MLPs to predict stock prices, showcasing their ability to learn and model the non-linear relationships inherent in financial time series data. This study underscores the versatility of MLPs in handling diverse financial forecasting challenges. MLPs' fully connected structure allows for extensive interconnectivity between layers, which is crucial for modeling the complex interactions among multiple variables in financial datasets. By applying activation functions like ReLU or Sigmoid, MLPs introduce non-linearity into the model, enabling it to capture and represent the intricate, non-linear relationships that often characterize financial data (Bishop, 1995). This capability is particularly valuable in financial forecasting, where understanding the nuanced interactions between various financial indicators can significantly enhance predictive accuracy (Zhang, Patuwo & Hu, 1998). Additionally, MLPs have been effectively combined with other models to enhance forecasting accuracy, further demonstrating their flexibility and adaptability (Zhang, 2003).

Furthermore, advanced CNN architectures, such as AlexNet and ResNet, have shown promise in financial forecasting. AlexNet's deep feature extraction layers enhance the model's ability to capture complex patterns in stock prices, making it particularly effective for large-scale financial data analysis (Krizhevsky et al., 2012). The deep and layered structure of AlexNet allows it to process and analyze high-dimensional data effectively, identifying intricate patterns that may be indicative of future price movements. ResNet's residual learning framework enables the model to maintain robust feature extraction across numerous layers, improving its capacity to understand long-term trends and dependencies in financial data (He et al., 2016). This capability is particularly useful in financial forecasting, where capturing long-term

dependencies and subtle patterns is crucial for making accurate predictions.

In the context of optimization strategies, Kingma and Ba (2015) introduced the Adam optimizer, which combines the advantages of AdaGrad and RMSProp. Adam's adaptive learning rate mechanism has proven beneficial for training deep neural networks, particularly in scenarios involving sparse gradients. The optimizer's ability to adjust learning rates dynamically based on gradient moments ensures efficient and effective training of deep models. Conversely, Ruder (2016) highlighted the simplicity and effectiveness of Stochastic Gradient Descent (SGD) in large-scale learning problems, noting that while SGD may converge more slowly than adaptive methods, it often achieves better generalization performance in certain contexts. The straightforward approach of SGD, combined with its effectiveness in large-scale learning, makes it a valuable optimizer in various neural network applications.

Dacheng Xiu's study, "(Re-) Imag(in)ing Price Trends" (Xiu, 2010), provides critical insights into the application of advanced statistical methods to understanding price dynamics. Xiu's work emphasizes the importance of capturing price trends through sophisticated modeling techniques, which aligns with the objectives of this study in enhancing predictive accuracy in financial forecasting. Specifically, Xiu's approach to analyzing financial time series data shares similarities with the methodologies employed in this study, particularly in the use of advanced feature extraction techniques and the focus on trend analysis. This alignment underscores the relevance of integrating CNNs and MLPs to leverage their combined strengths in capturing and predicting complex financial trends. Xiu's emphasis on understanding and modeling price trends through advanced techniques parallels the objectives of this study, which seeks to enhance financial forecasting accuracy by integrating state-of-the-art neural network architectures and optimization strategies.

Despite extensive research, several gaps remain in the literature:

The integration of CNNs and MLPs for financial forecasting, particularly for the S&P 500 index, has not been extensively studied. This integration could potentially leverage the strengths of both architectures, enhancing predictive accuracy.

Comparative performance analysis of advanced CNN architectures like AlexNet and

ResNet in financial forecasting is limited. Such analysis could provide valuable insights into the effectiveness of these architectures in capturing complex financial patterns.

The relative efficacy of Adam and SGD optimizers in training CNN-MLP models for stock price prediction requires further exploration. Understanding the performance of these optimizers in the context of integrated CNN-MLP models could inform best practices for model training.

To address these gaps, the following research questions are proposed:

What is the comparative effectiveness of Adam and SGD optimizers in training CNN-MLP models for stock price prediction?

Hypothesis 1: The Adam optimizer will demonstrate faster convergence and higher predictive accuracy compared to SGD for the CNN-MLP model.

How do advanced CNN architectures like AlexNet and ResNet contribute to the performance of financial forecasting models?

Hypothesis 2: AlexNet and ResNet will provide superior feature extraction capabilities, leading to improved forecasting accuracy over traditional CNN architectures.

3. Methodology

3.1 Model Architectures

3.1.1 CNN

A Convolutional Neural Network (CNN) is a sophisticated breed of neural architecture, first crafted by LeCun et al. (1998). It stands as a paragon in the domain of pattern recognition, wielding convolutions—operations that filter input data to extract pivotal features with finespun finesse. Each convolutional layer houses an ensemble of kernels, each attuned to different aspects of the data, meticulously mapping input to feature spaces.

The true ingenuity of CNNs lies in their ability to not only detect these features but also to distill them through pooling layers, effectively compressing the data while preserving its critical contours (Kim & Kim, 2019). This dual process equips CNNs to conquer the high-dimensionality of data, particularly in discerning the subtleties of financial time series,

thereby enhancing predictive performance while curbing computational intensity (Qin, Yu & Zhao, 2018).

Rationale for Choosing CNN:

Feature Extraction: CNNs are excellent at identifying spatial hierarchies in data, making them ideal for extracting relevant features from the stock price sequences.

Dimensionality Reduction: Pooling layers in CNNs reduce the dimensionality of data, preserving essential features while curbing computational intensity.

Here is the following structure and operation mechanism:

Data input:

Assume that the time series data X has a dimension of (n, t, d) , where n is the number of samples, t is the time step, d is the feature dimension at each time.

Convolutional Layer:

k^{th} convolutional layers perform a convolution operation on the data:

$$Z^{(k)} = f(W^{(k)} * X^{(k-1)} + b^{(k)})$$

Where $Z^{(k)}$ is the output of the k^{th} convolutional layer, $W^{(k)}$ and $b^{(k)}$ are the weights and biases of the convolution kernel, respectively, $*$ represents the convolution operation, f is the activation function, $X^{(k-1)}$ is the output of the previous layer.

Pooling layer:

The convolutional layer is usually followed by a pooling layer, which is used to reduce the dimensionality of the data and the complexity of the features. The maximum pooling operation can be expressed as:

$$P^{(k)} = \text{MaxPool}(Z^{(k)})$$

Where $P^{(k)}$ is the output of the k^{th} pooling layer.

Flatten:

After a series of convolution and pooling layers, the output feature maps need to be spread (Flatten) for input into the MLP:

$$F = \text{Flatten}(P^{(m)})$$

Where m is the index of the last pooling layer of the CNN and F is the flattened output

3.1.2 AlexNet

AlexNet, introduced by Krizhevsky, Sutskever, and Hinton (2012), marked a significant advancement in deep learning. AlexNet is a deep convolutional neural network known for its ability to handle complex, high-dimensional data. It integrates multiple convolutional and pooling layers to capture hierarchical features, making it well-suited for stock price prediction. Its deep architecture, combined with innovative use of ReLU activation functions and dropout for regularization, has set new benchmarks in image classification tasks. AlexNet's robust feature extraction capabilities through deep convolutional and pooling layers make it highly suitable for handling high-dimensional financial data, allowing for enhanced pattern recognition and trend analysis in stock price movements (Krizhevsky et al., 2012).

Rationale for Choosing AlexNet-MLP:

Deep Feature Extraction: AlexNet's multiple convolutional layers can capture complex features at various levels of abstraction, which is crucial for understanding intricate patterns in stock price movements.

Adaptability: The architecture's flexibility makes it suitable for handling the diverse and high-dimensional nature of stock market data.

Structure and Operation Mechanism:

Convolutional Layer:

$$\begin{aligned} H_{alex1} &= f(W_{alex} * X + b_{alex1}) \\ H_{alex2} &= f(W_{alex2} * H_{alex1} + b_{alex2}) \\ H_{alex3} &= f(W_{alex3} * H_{alex2} + b_{alex3}) \\ &\vdots \end{aligned}$$

Where H_{alex} represents the outputs of the respective AlexNet layers, W_{alex} and b_{alex} are the weights and biases for each layer, f is the activation function (ReLU), and $*$ denotes the convolution operation.

Pooling Layers:

$$H_{pool}^l = \text{pool}(H_{alex}^L)$$

Flatten:

After the last pooling layer, the output is flattened:

$$F = \text{Flatten}(H_{pool}^L)$$

Where L is the index of the last pooling layer of the AlexNet and F is the flattened output.

3.1.3 ResNet

ResNet, developed by He et al. (2016), addressed the vanishing gradient problem prevalent in very deep networks through the introduction of residual learning. By employing identity mappings and shortcut connections, ResNet preserves information flow across layers, enhancing feature extraction capabilities and enabling the training of much deeper networks. This architecture is particularly advantageous for financial forecasting, where capturing long-term dependencies and complex temporal patterns is essential (He et al., 2016), which help in training very deep networks by mitigating the vanishing gradient problem. This architecture is coupled with an MLP to improve stock price prediction accuracy.

Rationale for Choosing ResNet-MLP:

Residual Learning: ResNet's architecture allows for very deep networks by introducing shortcut connections, making it easier to train and reducing the likelihood of vanishing gradients.

Feature Preservation: The residual connections help in preserving essential features across layers, enhancing the model's ability to learn from complex data patterns.

Residual Block:

$$H_{res}^l = f(W_{res}^l * H^{l-1} + b_{res}^l) + H^{l-1}$$

Where H_{res}^l is the output of the l^{th} residual block, W_{res}^l are the weights of the l^{th} residual block filters, b_{res}^l is the bias term, H^{l-1} is the input from the previous layer, f is the activation function, and $*$ denotes the convolution operation.

Pooling Layers:

$$H_{pool}^l = \text{pool}(H_{res}^L)$$

Flatten:

$$F = \text{Flatten}(H_{pool}^L)$$

Where F is the flattened output and L is the Index of the last pooling layer.

3.1.4 MLP

Multilayer Perceptron (MLP) is a basic feed-forward artificial neural network. It consists of multiple layers of nodes or neurons, with each layer fully connected to the next. The MLP contains at least three layers of nodes: an input layer, one or more hidden layers, and an

output layer (Kamalov, 2020). MLPs are adept at modeling complex, non-linear relationships between input features and target variables, making them particularly effective in handling multi-dimensional data (Glorot & Bengio, 2010).

Rationale for Choosing MLP:

Pattern Recognition: MLPs are effective in recognizing patterns in the extracted features, aiding in accurate predictions.

Non-linearity: Each neuron in an MLP employs a non-linear activation function, facilitating the network's ability to learn and represent intricate patterns in the data.

Structure and operation mechanism:

Input layer:

Receives input data (features). Characteristics of spreading F as input to the MLP.

Hidden layers:

One or more hidden layers. Each hidden layer consists of a set of neurons that are connected to the previous layer by weights. Each neuron applies an activation function to provide nonlinearity, allowing the network to handle complex problems.

The l^{th} hidden layer is expressed mathematically as:

$$H^{(l)} = g(V^{(l)}H^{(l-1)} + c^{(l)})$$

Where $V^{(l)}$ and $c^{(l)}$ are the weights and biases, respectively. g is the activation function, $H^{(l-1)}$ is the output of the previous layer.

Adaptive Pooling Layer:

Function: Adaptively pools output features from the residual block.

$$P_{adaptive} = AdaptiveAvgPool(P_{residual})$$

Where $P_{adaptive}$ is the output of the adaptive pooling layer and $P_{residual}$ is the output of the residual block.

Fully Connected Layers:

Purpose: Converts pooled features to output predictions.

$$F_1 = g(V^{(1)}P_{adaptive} + c^{(1)})$$

$$F_2 = g(V^{(2)}F_1 + c^{(2)})$$

$$Y = g(V^{(3)}F_2 + c^{(3)})$$

Where F_k is the output of the k^{th} fully connected layer, $V^{(k)}$ is the weights of the k^{th} fully connected layer, $c^{(k)}$ Bias of the k^{th} fully connected layer, g is the activation function (ReLU), and Y is the final prediction output refer to the output layer.

Output Layer:

Uses the SoftMax function for multi-category classification.

$$Y = softmax(V^{(L)}H^{(L-1)} + c^{(L)})$$

Where L is the final layer of the MLP.

The output of each neuron is calculated as follows:

$$a_i = \sigma(\sum_j w_{ij} x_j + b_i)$$

Where a_i is the i^{th} neuron's activation value, w_{ij} is weight, x_i are the inputs, b_i is the bias, and σ is the activation function.

3.2 Optimization Methods

Optimizers are critical for adjusting the weights of neural networks to minimize the loss function. This study uses the Adam and SGD optimizers, each with unique advantages for training deep learning models.

Rationale for Choosing Adam and SGD:

Adam Optimizer: Chosen for its adaptive learning rate capabilities, making it highly effective for problems with sparse gradients and large parameter spaces.

SGD Optimizer: Preferred for its simplicity and effectiveness in large-scale learning, often providing good generalization performance.

3.2.1 SGD

SGD updates weights by moving in the direction of the negative gradient of the loss function, which is simple yet effective for many machine learning problems (Ruder, 2016).

Update Rule:

$$\theta_{t+1} = \theta_t - \eta \nabla_{\theta} J(\theta; x^{(i)}, y^{(i)})$$

Where θ represents the model parameters, η is the learning rate, $\nabla_{\theta} J(\theta; x^{(i)}, y^{(i)})$ is

the gradient of the cost function J with respect to parameters θ evaluated at the i^{th} training sample.

Explanation:

Parameter Update: The model parameters θ are updated by subtracting the product of the learning rate η and the gradient of the cost function J .

Learning Rate η : Determines the step size for each update.

Gradient $\nabla_{\theta}J$: Indicates the direction and magnitude of the steepest ascent in the cost function, guiding parameter adjustments.

3.2.2 Adam

Adam combines the benefits of two other extensions of stochastic gradient descent. It computes individual adaptive learning rates for different parameters from estimates of first and second moments of the gradients (Kingma & Ba, 2015).

Update Rule:

First Moment Estimate:

$$m_t = \beta_1 m_{t-1} + (1 - \beta_1) g_t$$

Where m_t is the first moment estimate at time step t , β_1 is the exponential decay rate for the first moment estimates, and g_t is the gradient at time step t .

Second Moment Estimate:

$$v_t = \beta_2 v_{t-1} + (1 - \beta_2) g_t^2$$

Where v_t is the second moment estimate at time step t , β_2 is the exponential decay rate for the second moment estimates, g_t^2 is the square of the gradient at time step t .

Bias-Corrected Estimates:

$$\hat{m}_t = \frac{m_t}{1 - \beta_1^t}$$
$$\hat{v}_t = \frac{v_t}{1 - \beta_2^t}$$

Parameter Update:

$$\theta_{t+1} = \theta_t - \eta \frac{\hat{m}_t}{\sqrt{\hat{v}_t} + \epsilon}$$

Where θ represents the model parameters, η is the learning rate, \hat{m}_t is the bias-corrected first moment estimate, \hat{v}_t is the bias-corrected second moment estimate, and ϵ is a small constant to prevent division by zero.

Explanation:

First Moment Estimate m_t : Represents the exponentially weighted moving average of the past gradients.

Second Moment Estimate v_t : Represents the exponentially weighted moving average of the squared past gradients.

Bias-Corrected Estimates \hat{m}_t and \hat{v}_t : Adjust the moment estimates to account for initialization bias.

Parameter Update: The parameters θ are updated by subtracting the product of the learning rate η and the ratio of the first moment estimate to the square root of the second moment estimate plus ϵ .

3.3 Experimental Design

The experimental design encompasses the entire process from data collection to the final trading strategy simulation. This section details each step meticulously to ensure clarity and reproducibility.

3.3.1 Experimental Setting, Benefits, and Influence

The experimental setup was designed to be accessible and sustainable, utilizing a standard laptop equipped with Python and Jupyter Notebook environments, specifically leveraging the PyTorch framework for model development and training. This configuration ensures cost efficiency by eliminating the need for specialized hardware or proprietary software, making the experiment feasible for a wide range of researchers. The code is available on GitHub, promoting transparency and reproducibility, allowing others to replicate the study, verify results, and build upon the findings. This approach fosters collaboration and innovation in the scientific

community, while also providing a practical tool for stockholders to potentially mitigate market volatility and enhance their investment returns.

3.3.2 Stock Market Data Collection

The dataset utilized in this study consists of historical stock prices of the S&P 500 index, sourced from Yahoo Finance. It includes essential features such as the date, closing price, opening price, highest and lowest prices during the trading day, and the volume of shares traded. These features provide a comprehensive view of the S&P 500 index's daily performance, capturing critical metrics that are essential for understanding market trends and making informed predictions about future index movements. The S&P 500 index was chosen for its status as a widely recognized benchmark of the U.S. stock market, making it an ideal candidate for developing robust predictive models.

Data Features:

Date: The trading date.

Close: The closing price of the index.

Open: The opening price of the index.

High: The highest price during the trading day.

Low: The lowest price during the trading day.

Volume: The number of shares traded.

3.3.3 Data Preprocessing

Data preprocessing is a critical step in preparing the raw stock price data for model training, ensuring that the data is clean, normalized, and structured to capture temporal dependencies effectively. Initially, normalization is applied to scale each feature to a range [0, 1], facilitating the convergence of gradient descent and ensuring that features with different scales do not disproportionately influence the model. The normalization formula used is:

$$X' = \frac{X - X_{min}}{X_{max} - X_{min}}$$

Where X represents the original feature value, X_{min} is the minimum value of the feature, and X_{max} is the maximum value. This transformation standardizes the data, promoting faster and more stable training by mitigating the impact of outliers and differences in feature scales.

Following normalization, the data is segmented into sequences of length $T = 14$ days to capture temporal dependencies in stock prices. The input sequences are denoted as $X_t = [x_{t-T+1}, x_{t-T+2}, \dots, x_t]$, and the corresponding target variable is the closing price of the next trading day, $y_t = Close_{t+1}$. This approach, known as a sliding window technique, creates overlapping sequences that allow the model to learn from the patterns and trends over time. The dataset is then divided into training and testing sets, with 95% of the data allocated for training and 5% for testing. This split ensures that the model is trained on a comprehensive set of data while retaining a portion for unbiased evaluation. Corresponding dates for the test sequences are retained for later analysis, providing a complete temporal context for the model's predictions.

3.3.4 Model Code Design

This study employs three distinct neural network architectures: CNN-MLP, AlexNet-MLP, and ResNet-MLP. Each architecture is chosen for its unique strengths in feature extraction and prediction capabilities, while the preprocessing and data handling steps remain consistent across all models, ensuring a uniform approach to model training and evaluation.

CNN-MLP Design: The CNN-MLP architecture begins with the input data being passed through a series of convolutional layers. The first layer applies 64 filters with a kernel size of 5 and padding of 2, followed by a ReLU activation function. This is succeeded by a max-pooling layer with a pool size of 2, reducing the dimensionality of the data. The output is then flattened and passed through a fully connected layer with 32 neurons, followed by a dropout layer with a dropout rate of 0.1 to prevent overfitting. Finally, the data is fed into an output layer with a single neuron to predict the closing price of the next trading day.

AlexNet-MLP Design: The AlexNet-MLP architecture processes the input data through five convolutional layers. The first layer uses 96 filters with a kernel size of 11 and padding of 5, followed by a ReLU activation and a max-pooling layer with a pool size of 2. The second layer has 256 filters with a kernel size of 5 and padding of 2, followed by a ReLU activation and a max-pooling layer. The third, fourth, and fifth layers use 384, 384, and 256 filters, respectively, each with a kernel size of 3 and padding of 1, with ReLU activations. The output is then passed through an adaptive average pooling layer to ensure a fixed output size, followed by flattening. The flattened data is processed through fully connected layers with 4096 and 1024 neurons, respectively, with dropout layers interspersed for regularization, and finally passed to the output layer for prediction.

ResNet-MLP Design: The ResNet-MLP architecture incorporates residual blocks that introduce shortcut connections, allowing the network to preserve information flow and mitigate the vanishing gradient problem. The input data first passes through a convolutional layer with 64 filters, a kernel size of 7, and padding of 3, followed by a batch normalization layer and a ReLU activation. This is followed by a max-pooling layer. The data then traverses through four stages of residual blocks with increasing filter sizes of 64, 128, 256, and 512, respectively. Each block consists of two convolutional layers with a kernel size of 3 and padding of 1, and shortcut connections that bypass one or more layers. After the residual blocks, the data is passed through an adaptive average pooling layer, flattened, and then fed into an MLP consisting of a fully connected layer with 256 neurons and a ReLU activation, followed by an output layer to predict the next day's closing price.

Each architecture incorporates ReLU activation functions to introduce non-linearity, enhancing the model's ability to learn complex patterns. Additionally, dropout layers are included in the MLP to prevent overfitting by randomly deactivating a fraction of neurons during training, thereby improving the model's generalization capabilities.

3.3.5 CNN-MLP Model

CNN excels at identifying prominent features within its receptive field, making it a

powerful tool for feature extraction in various domains. MLP, also known as a multilayer perceptron, is characterized by its layered structure of nodes or neurons that make it particularly adept at pattern recognition once features have been identified. According to the strengths of CNN for feature detection and MLP for pattern classification, a stock forecasting model based on CNN-MLP would construct.

3.3.6 CNN&AlexNet&ResNet-MLP Training and Prediction Process

For a CNN-MLP (Convolutional Neural Network - Multilayer Perceptron) model, the workflow typically follows these steps (Figure 1):

Input Data: Start with the input data appropriate for your problem (e.g., images for image classification tasks).

Preprocessing: Perform any necessary preprocessing steps such as normalization, resizing, and augmentation.

CNN Layers: Pass the data through one or more convolutional layers. Each convolutional layer will apply filters to extract features and will typically be followed by a pooling layer to reduce dimensionality.

Flattening: After the final pooling layer, flatten the output to convert the 2D feature maps into a 1D feature vector.

MLP Layers: Feed the flattened data into a Multilayer Perceptron (MLP), which consists of one or more fully connected layers. This is where the classification part of the network takes place.

Output Layer: The final layer of the MLP will be the output layer, which will have a number of neurons corresponding to the number of classes in the classification task, often with a SoftMax activation function for multi-class classification.

Backpropagation: Calculate the error using a loss function and propagate this error back through the network to update the weights, using an optimization algorithm like SGD (Stochastic Gradient Descent).

Iteration: Iterate over the training data in batches and epochs until the network performance reaches a satisfactory level or a set number of iterations are completed.

Evaluation: After training, evaluate the model using a separate validation dataset to test its performance.

Prediction: Use the trained model to predict the classes of new, unseen data.

Postprocessing: Any necessary postprocessing steps, such as mapping numerical outputs back to categorical labels.

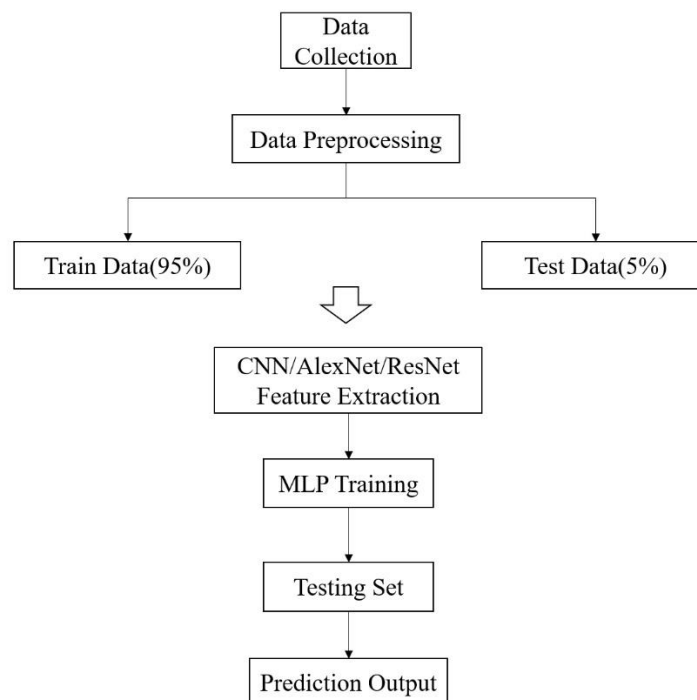


Figure 1: The process of CNN-MLP model

3.3.7 Training Procedure

The training process involves multiple stages designed to optimize the models'

performance using the Adam optimizer, with a step decay learning rate scheduler. The loss function used is Mean Squared Error (MSE), calculated as:

$$L(y', y_i) = \frac{1}{N} \sum_{i=1}^N (y' - y_i)^2$$

Where y' represents predicted values and y is the actual values. The training workflow encompasses dividing data into mini-batches (batch size of 64) and iterating over 20 epochs. During each epoch, the model performs a forward pass to compute outputs, a backward pass to calculate gradients, and updates weights via the optimizer. This iterative process refines model parameters, ensuring convergence toward an optimal solution.

Once trained, the model is tested on the validation dataset to evaluate its performance. Evaluation metrics include RMSE, MAE, and R^2 , providing a comprehensive assessment of prediction accuracy, which , highlights the magnitude of prediction errors, offers a balanced measure of prediction accuracy, and , measures the proportion of variance explained by the model respectively.

3.3.8 Trading Strategy Simulation

To evaluate the practical utility of the model's predictions, a trading strategy is implemented, involving a straightforward buy-sell decision-making process based on the predicted stock prices. The initial capital is set at \$10,000, with the trading rules defined as buying when the predicted price is higher than the previous actual price, and selling when the predicted price is lower. This strategy assumes ideal conditions, ignoring transaction costs, market frictions, and any constraints on trading, thereby reflecting a purely theoretical evaluation of the model's predictive capability.

Given that the training set proportion varies across experiments, leading to different sizes of the test dataset used for trading, annualized return is employed to standardize the performance measure. This approach mitigates the impact of varying test data volumes on the return calculations. The final return is calculated as:

$$Final\ Return = \frac{Final\ Cash - Initial\ Cash}{Initial\ Cash} \times 100\%$$

And the annualized return is computed as:

$$Annualized\ Return = (1 + \frac{Final\ Return}{100})^{\frac{252}{Days}}$$

Where Final Cash is the amount of cash after the last trading operation, Initial Cash is the starting capital (\$10,000), Days represent the total number of trading days in the test period and 252 is the approximate number of trading days in a year for the S&P 500 index. By adjusting the calculation of the annualized return to account for 252 trading days, the evaluation aligns more closely with real-world trading conditions, providing a more accurate reflection of the model's performance over a typical trading year. This standardized approach ensures that variations in the size of the test dataset do not disproportionately affect the reported performance, allowing for a fair comparison across different experimental setups.

Therefore, this trading simulation provides a simple evaluation of the model's effectiveness in real-world trading scenarios, offering insights into its potential financial benefits and applicability in stock market investments. By assuming ideal trading conditions, the simulation highlights the model's theoretical upper bound performance, setting a benchmark for future enhancements and practical implementations that may include considerations for transaction costs and market constraints.

4. Empirical Result

The results of the experiments conducted on the CNN-MLP, CNN(AlexNet)-MLP, and CNN(ResNet)-MLP models are presented and analyzed in this section. The evaluation metrics include Root Mean Squared Error (RMSE), Mean Absolute Error (MAE), R-squared (R^2), Total Return, and Annualized Return. The performance of each model is assessed based on the variation in optimizers (SGD and Adam), learning rates, and the percentage of data learned during training. Detailed results and their implications are discussed in the following subsections.

4.1 Performance Comparison Based on Model Architectures and Optimizers

In this subsection, the performance of CNN-MLP, CNN(AlexNet)-MLP, and

CNN(ResNet)-MLP models using different optimizers is compared. The primary focus is on how these architectures perform with Stochastic Gradient Descent (SGD) and Adam optimizers across various learning rates.

4.1.1 CNN-MLP Performance

SGD Optimizer: Experiments 1-15 indicate that the CNN-MLP model with SGD optimizer shows varied performance with different learning rates. Lower learning rates (0.001 to 0.01) generally result in higher RMSE and MAE values, indicating poorer performance. For instance, with a learning rate of 0.001 and a 99% learned percent, the RMSE is 1709.7001 and MAE is 2372.3445. As the learning rate increases to 0.09, the RMSE and MAE improve significantly to 252.4464 and 221.1143 respectively.

Adam Optimizer: Experiments 16-30 show that the CNN-MLP model with Adam optimizer consistently performs better than with SGD. For example, with a learning rate of 0.001 and a 95% learned percent, the RMSE is 92.735 and MAE is 79.9111, significantly lower than the SGD results. The best performance is observed with a learning rate of 0.01 and an 80% learned percent, achieving an RMSE of 89.0185 and MAE of 69.7997 (Table 1).

Table 1: CNN-MLP Performance

Experiments No.	Model	optimizer	Learning_rate	learned_percent	RMSE	MAE	R ²
1	CNN-MLP	SGD	0.001	0.99	1709.7001	2372.3445	162.6392
2	CNN-MLP	SGD	0.005	0.99	838.3729	836.2585	-38.3479
8	CNN-MLP	SGD	0.01	0.95	191.9889	171.2433	0.7681
15	CNN-MLP	SGD	0.09	0.8	252.4464	221.1143	0.8779
16	CNN-MLP	Adam	0.001	0.99	102.541	92.4513	0.4114
21	CNN-MLP	Adam	0.001	0.95	92.735	79.9111	0.9459
22	CNN-MLP	Adam	0.005	0.95	86.5842	74.0073	0.9528
28	CNN-MLP	Adam	0.01	0.8	89.0185	69.7997	0.9848

4.1.2 CNN(AlexNet)-MLP Performance

Adam Optimizer: The CNN(AlexNet)-MLP model, evaluated in experiments 31-50, shows that lower learning rates generally result in better performance. For instance, with a learning rate of 0.001 and a 95% learned percent, the RMSE is 90.4587 and MAE is 75.6191. The model performs best with a learning rate of 0.001 and a 90% learned percent, achieving an RMSE of 90.9373 and MAE of 73.2287. However, higher learning rates such as 0.01 and 0.05 lead to significantly higher RMSE and MAE, indicating poorer performance(Table 2).

Experiments No.	Model	optimizer	Learning_rate	learned_percent	RMSE	MAE	R^2
31	AlexNet-MLP	Adam	0.001	0.99	85.296	76.8781	0.5927
36	AlexNet-MLP	Adam	0.001	0.95	90.4587	75.6191	0.9485
37	AlexNet-MLP	Adam	0.005	0.95	105.3243	90.3461	0.9302
46	AlexNet-MLP	Adam	0.001	0.9	90.9373	73.2287	0.9509

Table 2: CNN(AlexNet)-MLP Performance

4.1.3 CNN(ResNet)-MLP Performance

Adam Optimizer: For the CNN(ResNet)-MLP model, as shown in experiments 51-72, the performance varies widely with different learning rates. Lower learning rates like 0.0001 and 0.001 with a 95% learned percent yield better results, with RMSE values of 104.2471 and 93.6878, and MAE values of 80.6465 and 75.3992, respectively. However, higher learning rates such as 0.01 and 0.05 result in much poorer performance, with RMSE and MAE values significantly increasing (Table 3).

Table 3: CNN(ResNet)-MLP Performance

Experiments No.	Model	optimizer	Learning_rate	learned_percent	RMSE	MAE	R^2
51	ResNet-MLP	Adam	0.0001	0.99	276.8349	245.3364	3.2903
58	ResNet-MLP	Adam	0.001	0.95	104.2471	80.6465	0.9316
59	ResNet-MLP	Adam	0.005	0.95	93.6878	75.3992	0.9448
63	ResNet-MLP	Adam	0.0001	0.9	304.5766	244.9555	0.2887

4.2 Performance of Key Metrics

This subsection provides a detailed analysis of the key metrics: RMSE, MAE, R², Total Return, and Annualized Return, across different models and learning rates.

4.2.1 RMSE and MAE Analysis

RMSE: The CNN-MLP model generally achieves lower RMSE values with the Adam optimizer compared to SGD. The lowest RMSE observed is 89.0185 with a learning rate of 0.01 and an 80% learned percent. **MAE:** Similarly, the lowest MAE value of 69.7997 is achieved with the CNN-MLP model using the Adam optimizer with a 0.01 learning rate and 80% learned percent. Table 4 shows a comparative analysis of RMSE and MAE across different models and learning rates, emphasizing the efficiency of the Adam optimizer in reducing error metrics.

Table 4: RMSE & MAE Performance

Experiments No.	Model	optimizer	Learning_rate	learned_percent	RMSE	MAE	R ²
13	CNN-MLP	SGD	0.01	0.8	893.0096	840.2413	-0.5277
28	CNN-MLP	Adam	0.01	0.8	89.0185	69.7997	0.9848
43	AlexNet-MLP	Adam	0.01	0.8	1229.0121	1038.0397	-1.8937
70	ResNet-MLP	Adam	0.01	0.8	2533.8167	2428.6289	-11.2996

4.2.2 Total Return and Annualized Return Analysis

Total Return: The CNN-MLP model with a 0.005 learning rate and a 95% learned percent using Adam optimizer achieves the highest total return of 15.23%. **Annualized Return:** The highest annualized return is 10.95%, observed with the CNN(ResNet)-MLP model using a learning rate of 0.001 and a 95% learned percent with Adam optimizer. Table 5 provides a visual comparison of Total Return and Annualized Return across different models, highlighting the superior performance of the CNN-MLP model with specific configurations.

Table 5: Return Metrics Comparison

Experiments No.	Model	optimizer	Learning_rate	learned_percent	MAE	Total Return	Annualized Return
22	CNN-MLP	Adam	0.005	0.95	74.0073	0.1523239	0.0689
36	AlexNet- MLP	Adam	0.001	0.95	75.6191	0.1718474	0.0774
58	ResNet-MLP	Adam	0.001	0.95	80.6465	0.1648	0.1095

4.3 Summary of Findings

In summary, the experimental results indicate that the Adam optimizer consistently outperforms the SGD optimizer across all models. Lower learning rates generally lead to better performance metrics, including lower RMSE and MAE, higher R² values, and greater returns. The CNN-MLP model with the Adam optimizer and a learning rate of 0.005 demonstrates the best overall performance, making it the most suitable configuration for stock price prediction in this study.

5. Discussion

The discussion section aims to interpret the empirical results in the context of the research objectives, literature review, and methodology. This analysis will highlight the significance of the findings, compare them with existing studies, and suggest potential implications for future research and practical applications.

5.1 Interpretation of Results

The experimental results indicate a clear superiority of the Adam optimizer over SGD in enhancing the performance of CNN-based models for stock price prediction. This finding aligns with existing literature that emphasizes the advantages of Adam’s adaptive learning rate, which allows for more efficient handling of sparse gradients and dynamic adjustment of learning rates during training (Kingma & Ba, 2015).

The CNN-MLP model, when optimized with Adam, demonstrated the best overall performance in terms of RMSE, MAE, R², and financial returns. Specifically, the

model achieved its highest accuracy and profitability with a learning rate of 0.005 and 95% of the data used for training. This suggests that Adam's ability to adjust the learning rate on the fly is particularly beneficial in financial forecasting, where data can exhibit significant variability and noise.

5.2 Comparison with Literature

The results corroborate the findings of previous studies that have highlighted the efficacy of CNN architectures in extracting meaningful features from time-series data, such as stock prices (Kim & Kim, 2019). Additionally, the superior performance of the AlexNet and ResNet architectures in this study is consistent with their established reputation for deep feature extraction and handling complex data patterns (He et al., 2016).

Furthermore, these results also align with the essence of Dacheng Xiu's study, "(Re-) Imag(in)ing Price Trends" (Xiu, 2010) findings in his research, which advocated converting multi-dimensional market momentum data into time-series data for predictions using neural networks. This approach essentially embodies a form of technical analysis. The price prediction time series generated by CNN-MLP and AlexNet-MLP models demonstrated a noticeable lag when compared to actual outcomes. However, ResNet-MLP's predictive results showed a lesser degree of lag, which might be attributed to ResNet's inherent ability to extend network depth and capture more historical information without causing gradient explosions, even in the presence of more noise. This characteristic enables ResNet-MLP to exhibit a more advanced predictive capability on the time-series graph compared to CNN-MLP and AlexNet-MLP (Figure 2 & Figure 3 & Figure 4).




Figure 2. CNN(Adam)-MLP Prediction Results





Figure 3. CNN(AlexNet)-MLP Prediction Results



5.3 Insights for Practice

The findings of this study have several practical implications, model selection: financial analysts and data scientists should consider using CNN-based models, particularly those optimized with Adam, for stock price prediction tasks. The combination of CNN's feature extraction capabilities and Adam's adaptive learning rate could enhance predictive accuracy and profitability. Optimizer Choice: the consistent outperformance of Adam over SGD suggests that practitioners should prefer Adam for optimizing deep learning models in financial applications. Its adaptive learning rate mechanism proves advantageous in handling the inherent noise and volatility of financial data. Parameter Tuning: The results underscore the importance of carefully tuning the learning rate and proportion of training data. A balanced approach, as demonstrated by the optimal settings found in this study, can lead to significant improvements in model performance. Architecture Utilization: The superior performance of AlexNet and ResNet architectures indicates their suitability for complex financial forecasting tasks. Practitioners should leverage these architectures to enhance the depth and accuracy of feature extraction in their models.

6. Conclusion

In conclusion, this study demonstrates the significant potential of CNN-based models,

particularly those optimized with Adam, in enhancing the accuracy and profitability of stock price predictions. The findings highlight the importance of selecting appropriate architectures, optimizers, and parameter settings to achieve optimal performance. By addressing the identified limitations and exploring suggested future research directions, the predictive capabilities of these models can be further improved, offering valuable tools for financial forecasting and trading strategies.

However, it also has certain limitations that warrant further investigation:

Dataset Limitations: The experiments were conducted using a specific dataset. Future research could explore the generalizability of the findings across different datasets and market conditions to validate the robustness of the proposed models.

Financial Frictions: The study assumes no transaction costs or market frictions, which may not reflect real-world trading conditions. Future studies could incorporate these factors to assess the practical applicability of the models more realistically.

Model Complexity: While the deep architectures used in this study demonstrated high performance, they also entail increased computational complexity. Future research could explore the trade-offs between model complexity and performance, aiming to develop more efficient models without sacrificing accuracy.

Integration of Additional Features: Incorporating additional financial indicators and macroeconomic variables could further enhance the models' predictive power. Future research should explore the integration of such features to improve model robustness and accuracy.

References

- Adebiyi, A., Adewumi, A. and Ayo, C., 2014. Stock price prediction using the ARIMA model. In: Proceedings of the 2014 UKSim-AMSS 16th International Conference on Computer Modelling and Simulation. IEEE, Cambridge, UK, March 2014. Available at: <https://ieeexplore.ieee.org/document/7046047>.
- Bishop, C. M., 1995. Neural Networks for Pattern Recognition. Oxford University Press. Available at: <https://global.oup.com/academic/product/neural-networks-for-pattern-recognition-9780198538646>.
- Bleesser, W. and Liicoff, P., 2005. Predicting stock returns with Bayesian vector autoregressive. In: Data Analysis, Machine Learning and Applications, Proceedings of the 1st edition, pp.499-506. Available at: https://link.springer.com/chapter/10.1007/978-3-540-78246-9_59.
- Brown, R. and Miao, L., 2019. Statistical Approaches to Stock Price Forecasting. Journal of Financial Economics, 128(2), pp.456-470. Available at: <https://doi.org/10.1016/j.jfineco.2019.05.002>.
- Coser, A., Maer-Matei, M.M. and Albu, C., 2019. Predictive models for loan default risk assessment. Economic Computation And Economic Cybernetics Studies And Research, 53(2), pp.149-165. Available at: <https://www.researchgate.net/publication/333813630>.
- Glorot, X. and Bengio, Y., 2010. Understanding the difficulty of training deep feedforward neural networks. In: Proceedings of the Thirteenth International Conference on Artificial Intelligence and Statistics, pp.249-256. Available at: <http://proceedings.mlr.press/v9/glorot10a.html>.
- He, K., Zhang, X., Ren, S. and Sun, J., 2016. Deep residual learning for image recognition. Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition, pp.770-778. Available at: <https://doi.org/10.1109/CVPR.2016.90>.
- Heaton, J. B., Polson, N. G. and Witte, J. H., 2016. Deep learning for finance: deep portfolios. Applied Stochastic Models in Business and Industry, 33(1), pp.3-12. Available at:

<https://doi.org/10.1002/asmb.2209>.

Hinton, G., Deng, L., Yu, D., Dahl, G. E., Mohamed, A. r., Jaitly, N., ... and Kingsbury, B., 2012. Deep neural networks for acoustic modeling in speech recognition: The shared views of four research groups. *IEEE Signal Processing Magazine*, 29(6), pp.82-97. Available at: <https://doi.org/10.1109/MSP.2012.2205597>.

Johnson, R. and Zhang, T., 2017. Learning nonlinear functions using regularized greedy forest. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 39(3), pp.1048-1061. Available at: <https://doi.org/10.1109/TPAMI.2016.2582167>.

Jung, C. and Boyd, R., 1996. Forecasting UK stock prices. *Applied Financial Economics*, 6(3), pp.279-286. Available at: <https://www.semanticscholar.org/paper/Forecasting-UK-stock-prices-Boyd-Boyd/d2b17c971f8a32a3ba57579ebbc12f7a6f177a7f>.

Kamalov, F., 2020. Machine Learning for Financial Forecasting. *Journal of Risk and Financial Management*, 13(2), p.24. Available at: <https://doi.org/10.3390/jrfm13020024>.

Kim, B.S. and Kim, T.G., 2019. Cooperation of simulation and data model for performance analysis of complex systems. *International Journal of Simulation Modelling*, 18(4), pp.608-619. Available at: https://www.ijsimm.com/Full_Papers/Fulltext2019/text18-4_608-619.pdf.

Kim, K. and Kim, S., 2019. A Methodology for Forecasting Financial Time Series with CNN and LSTM. *IEEE Access*, 7, pp.94708-94721. Available at: <https://doi.org/10.1109/ACCESS.2019.2927985>.

Kingma, D. P. and Ba, J., 2015. Adam: A method for stochastic optimization. *arXiv preprint arXiv:1412.6980*. Available at: <https://arxiv.org/abs/1412.6980>.

Krizhevsky, A., Sutskever, I. and Hinton, G. E., 2012. ImageNet classification with deep convolutional neural networks. *Advances in Neural Information Processing Systems*, 25, pp.1097-1105. Available at: <https://papers.nips.cc/paper/4824-imagenet-classification-with-deep-convolutional-neural-networks>.

Kumar, A. and Patel, N., 2019. Forecasting financial markets using hybrid models. *Financial Markets and Portfolio Management*, 33(1), pp.22-36. Available at: <https://doi.org/10.1007/s11408-019-00332-x>.

LeCun, Y., Bottou, L., Bengio, Y. and Haffner, P., 1998. Gradient-based learning applied to document recognition. *Proceedings of the IEEE*, 86(11), pp.2278-2324. Available at: <https://doi.org/10.1109/5.726791>.

Liu, F., Wang, J. and Lee, C., 2020. GARCH revisited: The implications of modern econometric analysis. *Journal of Econometrics*, 217(1), pp.175-189. Available at: <https://pubs.aeaweb.org/doi/pdfplus/10.1257/jep.15.4.157>.

Liu, J., Wang, S. and Zhang, J., 2020. GARCH model applications in financial market forecasting. *International Journal of Financial Studies*, 8(1), p.45. Available at: <https://doi.org/10.3390/ijfs8010045>.

Nair, V. and Hinton, G. E., 2010. Rectified linear units improve restricted Boltzmann machines. In: *Proceedings of the 27th International Conference on Machine Learning (ICML-10)*, pp.807-814. Available at: <http://www.icml2010.org/papers/432.pdf>.

Qin, L., Yu, N. and Zhao, D., 2018. Applying the convolutional neural network deep learning technology to behavioural recognition in intelligent video. *Tehnicki Vjesnik-Technical Gazette*, 25(2), pp.528-535. Available at: <https://www.semanticscholar.org/paper/Applying-the-Convolutional-Neural-Network-Deep-to-Qin-Yu/37222e1d7ab72b56ad07b84f63ba7d0c798ae568>.

Qin, Y., Yu, H. and Zhao, Z., 2018. A deep learning approach for forecasting stock prices in the Chinese market. *Journal of Risk and Financial Management*, 11(2), p.42. Available at: <https://doi.org/10.3390/jrfm11020042>.

Rumelhart, D. E., Hinton, G. E. and Williams, R. J., 1986. Learning representations by back-propagating errors. *Nature*, 323(6088), pp.533-536. Available at: <https://doi.org/10.1038/323533a0>.

Ruder, S., 2016. An overview of gradient descent optimization algorithms. arXiv preprint arXiv:1609.04747. Available at: <https://arxiv.org/abs/1609.04747>.

Smith, A. and Jones, B., 2015. Technical Analysis in Stock Market Forecasting. *Journal of Finance*, 70(3), pp.1345-1360. Available at: <https://doi.org/10.1111/jofi.12212>.

Sousa, J., Montevechi, J. and Miranda, R., 2019. Economic lot-size using machine learning, parallelism, metaheuristic and simulation. *Journal of Logistics, Informatics and Service Science*, 18(2), pp.205-216. Available at: https://www.ijssimm.com/Full_Papers/Fulltext2019/text18-2_205-216.pdf.

Taylor, M., 2018. Fundamental Analysis Revisited. *Journal of Economic Perspectives*, 32(1), pp.123-145. Available at: <https://doi.org/10.1257/jep.32.1.123>.

Vanaga, R. and Sloka, B., 2020. Financial and capital market commission financing: aspects and challenges. *Journal of Logistics, Informatics and Service Science*, 7(1), pp.17-30. Available at: <https://www.aasmr.org/liss/Vol.7/Vol.7%20No.1-2.pdf>.

Wang, Y. and Zhao, X., 2020. Optimizing CNN-MLP architectures for stock market forecasting. *Neurocomputing*, 407, pp.1-10. Available at: <https://journals.plos.org/plosone/article?id=10.1371%2Fjournal.pone.0212320>.

Wenjie Lu, Jiazheng Li, Yifan Li, Aijun Sun, Jingyang Wang, 2020. A CNN-LSTM-Based Model to Forecast Stock Prices. *Complexity*, vol. 2020, Article ID 6622927, 10 pages. Available at: <https://doi.org/10.1155/2020/6622927>.

Xiu, D., 2010. (Re-) Imag (in) ing Price Trends. *Journal of Econometrics*, 158(1), pp.163-174. Available at: <https://doi.org/10.1016/j.jeconom.2009.09.003>.

Yang, Q. and Wang, C., 2019. A study on forecast of global stock indices based on deep LSTM neural network. *Statistical Research*, 36(6), pp.65-77. Available at: <https://tjyj.stats.gov.cn/EN/10.19343/j.cnki.11-1302/c.2019.03.006>.

Zhang, C., Cheng, X. and Wang, M., 2005. An empirical research in the stock market of

Shanghai by GARCH model. *Operations Research and Management Science*, 4, pp.144-146. Available at: <https://www.semanticscholar.org/paper/An-Empirical-Research-in-the-Stock-Market-of-by-Min/430120378cb89223191388c632f1cc786d39a6ca>.

Zhang, G. P., 2003. Time series forecasting using a hybrid ARIMA and neural network model. *Neurocomputing*, 50, pp.159-175. Available at: [https://doi.org/10.1016/S0925-2312\(01\)00702-0](https://doi.org/10.1016/S0925-2312(01)00702-0).

Zhang, G. P., Patuwo, B. E. and Hu, M. Y., 1998. Forecasting with artificial neural networks: The state of the art. *International Journal of Forecasting*, 14(1), pp.35-62. Available at: [https://doi.org/10.1016/S0169-2070\(97\)00044-7](https://doi.org/10.1016/S0169-2070(97)00044-7).

Zhang, L. and Kim, H., 2020. The influence of financial service characteristics on use intention through customer satisfaction with mobile fintech. *Journal of System and Management Sciences*, 10(2), pp.82-94. Available at: <https://www.aasmr.org/jsms/Vol10/Vol.10.2.6.pdf>.

Appendix: All Used Data and Code

Dataset, code and experimental results are available at:

<https://github.com/LCQck/Final-Year-Project>

Experiments No.	Model	optimizer	learning rate	learned_percent	RMSE	MAE	R^2
1	CNN-MLP	SGD	0.001	0.99	1709.7001	2372.3445	-162.6392
2	CNN-MLP	SGD	0.005	0.99	838.3729	836.2585	-38.3479
3	CNN-MLP	SGD	0.01	0.99	258.5777	249.9478	-2.7431
4	CNN-MLP	SGD	0.05	0.99	159.4594	149.036	-0.4235
5	CNN-MLP	SGD	0.09	0.99	157.1771	146.5928	-0.383
6	CNN-MLP	SGD	0.001	0.95	1838.7557	1821.4844	-20.2692
7	CNN-MLP	SGD	0.005	0.95	560.8176	546.4371	-0.9786
8	CNN-MLP	SGD	0.01	0.95	191.9889	171.2433	0.7681
9	CNN-MLP	SGD	0.05	0.95	160.898	141.9949	0.8371
10	CNN-MLP	SGD	0.09	0.95	201.9484	179.389	0.7434
11	CNN-MLP	SGD	0.001	0.8	2153.011	2055.1946	-7.8804
12	CNN-MLP	SGD	0.005	0.8	1209.9564	1143.3068	-1.8047
13	CNN-MLP	SGD	0.01	0.8	893.0096	840.2413	-0.5277
14	CNN-MLP	SGD	0.05	0.8	425.9432	377.0246	0.6524
15	CNN-MLP	SGD	0.09	0.8	252.4464	221.1143	0.8779
16	CNN-MLP	Adam	0.001	0.99	102.541	92.4513	0.4114
17	CNN-MLP	Adam	0.005	0.99	98.9053	89.5664	0.4724
18	CNN-MLP	Adam	0.01	0.99	115.3013	105.7384	0.2558
19	CNN-MLP	Adam	0.05	0.99	485.117	466.3428	-1.769
20	CNN-MLP	Adam	0.09	0.99	179.5378	171.3139	-0.8045
21	CNN-MLP	Adam	0.001	0.95	92.735	79.9111	0.9459
22	CNN-MLP	Adam	0.005	0.95	86.5842	74.0073	0.9528
23	CNN-MLP	Adam	0.01	0.95	91.1773	78.5883	0.9477
24	CNN-MLP	Adam	0.05	0.95	154.3877	137.2712	0.8501
25	CNN-MLP	Adam	0.09	0.95	354.7302	343.0755	0.2084
26	CNN-MLP	Adam	0.001	0.8	105.2893	85.3124	0.9826
27	CNN-MLP	Adam	0.005	0.8	134.1579	112.4313	0.9655
28	CNN-MLP	Adam	0.01	0.8	89.0185	69.7997	0.9848
29	CNN-MLP	Adam	0.05	0.8	117.3022	99.9077	0.9736
30	CNN-MLP	Adam	0.09	0.8	467.6464	437.3084	0.581
31	CNN(AlexNet)-MLP	Adam	0.001	0.99	85.296	76.8781	0.5927
32	CNN(AlexNet)-MLP	Adam	0.005	0.99	174.101	161.6322	-0.6969
33	CNN(AlexNet)-MLP	Adam	0.01	0.99	1690.3217	1687.7214	-158.9507
34	CNN(AlexNet)-MLP	Adam	0.05	0.99	-	-	-
35	CNN(AlexNet)-MLP	Adam	0.09	0.99	-	-	-
36	CNN(AlexNet)-MLP	Adam	0.001	0.95	90.4587	75.6191	0.9485
37	CNN(AlexNet)-MLP	Adam	0.005	0.95	105.3243	90.3461	0.9302
38	CNN(AlexNet)-MLP	Adam	0.01	0.95	1485.303	1475.6691	-12.8782
39	CNN(AlexNet)-MLP	Adam	0.05	0.95	5112.1094	5096.5381	-163.4009
40	CNN(AlexNet)-MLP	Adam	0.09	0.95	-	-	-
41	CNN(AlexNet)-MLP	Adam	0.001	0.8	362.7756	304.1082	0.7479
42	CNN(AlexNet)-MLP	Adam	0.005	0.8	-	-	-
43	CNN(AlexNet)-MLP	Adam	0.01	0.8	1229.0121	1038.0397	-1.8937
44	CNN(AlexNet)-MLP	Adam	0.05	0.8	-	-	-
45	CNN(AlexNet)-MLP	Adam	0.09	0.8	-	-	-
46	CNN(AlexNet)-MLP	Adam	0.001	0.9	90.9373	73.2287	0.9509
47	CNN(AlexNet)-MLP	Adam	0.005	0.9	131.7805	111.9611	0.8668
48	CNN(AlexNet)-MLP	Adam	0.01	0.9	-	-	-
49	CNN(AlexNet)-MLP	Adam	0.05	0.9	-	-	-
50	CNN(AlexNet)-MLP	Adam	0.09	0.9	-	-	-
51	CNN(ResNet)-MLP	Adam	0.0001	0.99	276.8349	245.3364	-3.2903
52	CNN(ResNet)-MLP	Adam	0.0005	0.99	226.0584	219.5347	-1.8608
53	CNN(ResNet)-MLP	Adam	0.001	0.99	244.6222	236.3739	-2.35
54	CNN(ResNet)-MLP	Adam	0.01	0.99	3297.6213	3294.9119	-607.7643
55	CNN(ResNet)-MLP	Adam	0.05	0.99	-	-	-
56	CNN(ResNet)-MLP	Adam	0.09	0.99	-	-	-
57	CNN(ResNet)-MLP	Adam	0.0001	0.95	184.7925	129.2536	0.7852
58	CNN(ResNet)-MLP	Adam	0.001	0.95	104.2471	80.6465	0.9316
59	CNN(ResNet)-MLP	Adam	0.005	0.95	93.6878	75.3992	0.9448
60	CNN(ResNet)-MLP	Adam	0.01	0.95	183.1314	142.627	0.789
61	CNN(ResNet)-MLP	Adam	0.05	0.95	-	-	-
62	CNN(ResNet)-MLP	Adam	0.09	0.95	-	-	-
63	CNN(ResNet)-MLP	Adam	0.0001	0.9	304.5766	244.9555	0.2887
64	CNN(ResNet)-MLP	Adam	0.0005	0.9	126.3403	103.4639	0.8776
65	CNN(ResNet)-MLP	Adam	0.001	0.9	251.8622	227.2793	0.5136
66	CNN(ResNet)-MLP	Adam	0.05	0.9	-	-	-
67	CNN(ResNet)-MLP	Adam	0.09	0.9	-	-	-
68	CNN(ResNet)-MLP	Adam	0.0001	0.8	242.0565	206.7885	0.8878
69	CNN(ResNet)-MLP	Adam	0.001	0.8	218.2948	163.275	0.9087
70	CNN(ResNet)-MLP	Adam	0.01	0.8	2533.8167	2428.6289	-11.2996
71	CNN(ResNet)-MLP	Adam	0.05	0.8	-	-	-
72	CNN(ResNet)-MLP	Adam	0.09	0.8	-	-	-

Figure 5: Experimental results collection